

Capstone Project Report

James Cook University Cairns

Hunter Kruger-Ilingworth (14198489)

August 17, 2025

Contents

1	Introduction	2
2	Dataset	3
3	Model Structure	4
3.1	Model Structure	4
3.2	Hyperparameter Tuning with Hyperband	4
4	Model Deployment	4
5	Transfer Learning	6
6	Model Comparison And Evaluation	6
7	AWS Sagemaker Information and Discussion	7
8	Final Model Discussion and Conclusion	7
9	References	9
A	Appendix	10

List of Figures

1	AI generated images [1]	2
2	Screenshot of S3 Bucket in AWS Console	3
3	Endpoint page screenshot in Sagemaker, as evidence of deployment	6
4	Confusion Matrix of main model and transfer learning model	7
5	Project Cost summary	8

List of Tables

1	Dataset features and their descriptions.	3
2	Counts of each class label in the training and testing sets	3
3	Tunable Hyperparameters and Final Chosen Values	4
4	Model Comparison	6

Introduction

"AI as a creative tool has improved leaps and bounds since its early incarnations. Generating still images that are interesting, sophisticated and photorealistic is now an easy process that can be done by anybody with an interest, some patience and determination" [2]. The prevalence of AI among the general public has led to extensive dialogue about their ethical implications in producing artwork, or for manufacturing misinformation. One such example of the use of AI is tools like *DALL-E* and *Stable Diffusion*, which generate images from text prompts. In the current digital age, we cannot rely on the transparency of the source of an image, so there exists a need to be able to identify whether an image has been generated by an AI, or a real artist. Figure 1 shows some examples of AI generated images, which are becoming increasingly difficult to distinguish from real images as the technology improves, underscoring the need for reliable detection methods.



Fig. 1. AI generated images [1]

[1] evaluates both human and AI capabilities in detecting fake images, showing that humans are often deceived by advanced image generation models, while AI-based detection algorithms outperform humans but still misclassify 13% of images. The study introduction of the Fake2M dataset and new benchmarks (HPBench and MPBench) aims to drive further research and improve the reliability of AI-generated content detection.

[3] presents a computer vision approach to distinguishing AI-generated images from real ones, utilising a synthetic dataset created with Latent Diffusion, classification via Convolutional Neural Networks, and interpretability through Grad-CAM. Achieving nearly 93% accuracy, the study also introduces the CIFAKE dataset, a large collection of real and synthetic images, to support further research on the detection of AI-generated imagery.

It is therefore evident that past research has shown that it is possible to classify images as AI-generated or not, with a high degree of accuracy. This report will explore the process of training a neural network to classify images as either AI-generated or not, and then deploying this model to AWS SageMaker for inference, and evaluating it against comparable models. This is done so that people can distinguish between AI-generated and real images, ultimately as a tool to identify the spread of misinformation and other harms.

Dataset

The dataset used for this project was a competition dataset from Hugging Face, held in 2023 [4]. The dataset consists of 62,060 images, and is 2.37GB in size, being pre-split into training and testing sets, as summarised in tables 1 and 2, where it can be seen that the testing set has the class labels withheld due to the competition setting, restricting this analysis to the 18,618 training images, which we can sub-divide and validate with known labels.

Feature	Description
id	Index filename 34.jpg
image	The Image (512x512)
label	Binary class label

Table 1: Dataset features and their descriptions.

Class Label	Train Count	Test Count
AI (1)	10,330 (55.5%)	NA
Not AI (0)	8,288 (45.5%)	NA
Total	18,618	43,442

Table 2: Counts of each class label in the training and testing sets

Listing 2 shows the code used to load the dataset, and split it into training, validation, and test sets. It can be seen that the splitting method for the data was to - holdout 500 images for final testing (only 500 was chosen due to a bandwidth issue on the AWS endpoint, further elaborated on in Section 4), - split the remaining training data into three sets, the small train, train and test set, being 10%, 70%, and 20% of the remaining data respectively. - to ensure class balance between these sets, random oversampling was employed. This was because table 2 shows that if a model were to guess AI all the time, it could get an accuracy of 55.5%

Listing 2, in the appendix shows the code used to load the dataset from HuggingFace. It can then be seen that the splitting scheme described above is done using hugging faces `train_test_split` method. The HuggingFace dataset object is converted to numpy arrays. Then Imblearn's `RandomUnderSampler` is used to ensure class balance, as to not bias the majority class.

This data was converted to npz files, due to their efficient storage and loading capabilities. After conversion, the data was uploaded to an S3 bucket for easy access during model training and evaluation. Uploading to S3 buckets was critical, as it meant that there was no longer a reliance on the RAM allocation of the sagemaker notebook environment - it could be loaded in from the S3 bucket in batches, which has the benefit of being a more scalable solution.

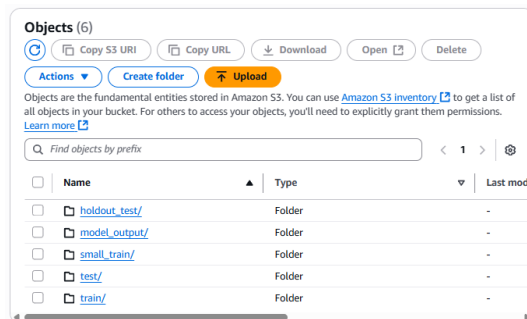


Fig. 2. Screenshot of S3 Bucket in AWS Console

Model Structure

- detail the structure of the model, eg. the training was called in a notebook (main.ipynb) but the training and the model was defined in discrete python files of their own in accordance with modularity best practices.

Model Structure

The model is a convolutional neural network (CNN) designed for binary image classification. It accepts full-size images of shape $(512, 512, 3)$.

Input $(512 \times 512 \times 3) \rightarrow \text{Rescaling}(1/255) \rightarrow \text{Conv2D}(f_1, 3 \times 3, \text{ReLU}) \rightarrow \text{Pool}(\text{max/avg}) \rightarrow \text{Conv2D}(f_2, 3 \times 3, \text{ReLU}) \rightarrow \text{Pool}(\text{max/avg}) \rightarrow \text{Global}(\text{max/avg}) \text{ pooling} \rightarrow \text{Flatten} \rightarrow \text{Dense}(d, \text{ReLU}) \rightarrow \text{optional Dropout}(p) \rightarrow \text{Dense}(1, \text{sigmoid})$. Trained with binary cross-entropy and adam/adagrad.

Adam was chosen as one of the considered optimisers, as it is known to converge rapidly, and rectifies vanishing learning rate and high variance, therefore being the most popular optimiser [5]

Adagrad was chosen as the second potential optimiser, as the learning rate would not need manual tuning, and is known to perform better than alternatives like SGD, MBGD, and primitive momentum based optimisers. Though it should be noted that it has the weakness of a constantly decreasing learning rate, resulting in slower convergence [5]

mention that the most optimal loss function for binary classification task is binary cross entropy.

mention the fact that i trained on 3 epochs

Hyperparameter Tuning with Hyperband

Hyperparameter tuning was performed to maximise the models performance through iteration of most of its parameters.

Table 3 shows a list of the parameters being tuned in the above model. It is evident that this is a very large parameter space, and it is therefore not feasible to find the best possible hyperparameter.

The Hyperband algorithm is used for hyperparameter optimization. Hyperband eliminates poorly performing hyperparameter combinations and focuses on promising ones, thus saving computational time compared to other methods like Bayesian optimization and grid search.

Table 3: Tunable Hyperparameters and Final Chosen Values

Hyperparameter	Range/Choices	Final Value
learning rate	1×10^{-4} to 1×10^{-2} (log scale)	0.00149
dropout-rate	0.0 to 0.5 (used if use-dropout=true)	0.284
conv1-filters (f_1)	16 to 128	24
conv2-filters (f_2)	32 to 256	107
dense-units (d)	64 to 512	254
pooling	max, avg	max
use-dropout	true, false	true
optimizer	adam, adagrad	adam

Model Deployment

The best hyperparameters found on the smaller training set were then extracted, and trained on the larger dataset. This

Listing 1. Hyperparameter tuning code extracts from `main.ipynb` and `src/train.py`

```

1 # main.ipynb
2
3 estimator = tf(
4     entry_point="train.py",
5     source_dir="src",
6     role=role,
7     use_spot_instances=True, # save money
8     instance_type="ml.c5.2xlarge",
9     instance_count=1,
10    framework_version="2.14",
11    py_version="py310",
12    hyperparameters={
13        "epochs": 3,
14        "height": 512,
15        "width": 512,
16        "channels": 3
17    },
18    output_path=s3_output_location
19 )
20
21 hyperparameter_ranges = {
22     "learning-rate": ContinuousParameter(1e-4, 1e-2, scaling_type="Logarithmic"),
23     "dropout-rate": ContinuousParameter(0.0, 0.5),
24     "batch-size": IntegerParameter(4, 8),
25     "conv1-filters": IntegerParameter(16, 128),
26     "conv2-filters": IntegerParameter(32, 256),
27     "dense-units": IntegerParameter(64, 512),
28     "pooling": CategoricalParameter(["max", "avg"]),
29     "use-dropout": CategoricalParameter(["true", "false"]),
30     "optimizer": CategoricalParameter(["adam", "adagrad"]),
31 }
32
33 metric_definitions = [
34     {"Name": "val_auc", "Regex": "val_auc: ([0-9\\.]+)"},
35     {"Name": "val_f1", "Regex": "val_f1: ([0-9\\.]+)"},
36     {"Name": "val_precision", "Regex": "val_precision: ([0-9\\.]+)"},
37     {"Name": "val_recall", "Regex": "val_recall: ([0-9\\.]+)"},
38     {"Name": "val_accuracy", "Regex": "val_accuracy: ([0-9\\.]+)"},
39 ]
40
41 tuner = HyperparameterTuner(
42     estimator=estimator,
43     objective_metric_name="val_f1",
44     strategy='Hyperband',
45     hyperparameter_ranges=hyperparameter_ranges,
46     metric_definitions=metric_definitions,
47     max_parallel_jobs=5,
48     objective_type="Maximize",
49     # early_stopping_type="Auto", # not supported for hyperband strategy, since it gets rid of unpromising trials itself
50     max_jobs=20,
51     base_tuning_job_name="ph-17",
52 )
53
54 tuner.fit({
55     "train": small_train_input,
56     "test": test_input,
57 })
58
59
60 # src/train.py
61 model = build_model(
62     input_shape=(args.height, args.width, args.channels),
63     conv1_filters=args.conv1_filters,
64     conv2_filters=args.conv2_filters,
65     dense_units=args.dense_units,
66     use_dropout=args.use_dropout,
67     dropout_rate=args.dropout_rate,
68     pooling=args.pooling,
69 )
70 if args.optimizer == "adagrad":
71     optimizer_choice = tf.keras.optimizers.Adagrad(learning_rate=args.learning_rate)
72 else:
73     optimizer_choice = tf.keras.optimizers.Adam(learning_rate=args.learning_rate)
74
75 model.compile(
76     optimizer=optimizer_choice,
77     loss="binary_crossentropy",
78     metrics=[
79         tf.keras.metrics.BinaryAccuracy(name="binary_accuracy"),
80         tf.keras.metrics.AUC(name="auc"),
81         tf.keras.metrics.Precision(name="precision"),
82         tf.keras.metrics.Recall(name="recall")]
83 )

```

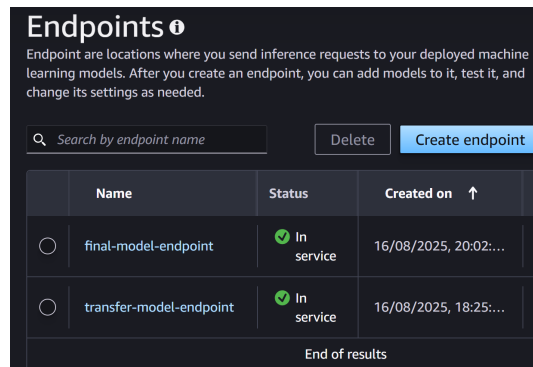


Fig. 3. Endpoint page screenshot in Sagemaker, as evidence of deployment

Transfer Learning

The list of the possible CNN structures available for use for transfer learning can be seen in [6]. Of these, EfficientNetV2 was selected as the most suitable architecture due to its demonstrated balance between accuracy and computational efficiency on general image classification tasks. This makes it well-suited for the diverse and varied images present in the aiornot dataset [6, 7].

Using an endpoint, the model was then evaluated on its precision, recall, f1 score, and accuracy. Figure 3 shows the endpoint configuration in Sagemaker. This allowed for evaluation of the models performance, whilst not being limited by the RAM limitations of the notebook environment in Sagemaker studio.

Deploying the transfer learning model, again using a sagemaker endpoint allowed a comparison between both models on the holdout set.

Model Comparison And Evaluation

In this section, we compare the performance of the different models trained on the dataset. Evaluation metrics were calculated by invoking the model endpoints on the holdout set, which was not used during training or validation. This ensures that the models' performance is on unseen data, better modelling a general use case.

The endpoints were deployed as real time endpoints, operating on a server and interacting through the user with a HTTP request. While this does mean they can be accessed at any time by any user who has the appropriate permissions, it also means that the network requests are limited to **X Mb**, which equates to only 4 500x500 rgb images at a time, making the evaluation process slower. It was for this reason that the holdout set was limited to only 500 images.

Table 4: Model Comparison

Model	Accuracy	Precision	Recall	F1 Score
Main Model	0.854	0.833	0.906	0.868
Transfer Model	0.856	0.875	0.849	0.862
Difference	0.002	0.042	0.057	0.006

The results of the model evaluation are summarized in Table 4. The table highlights the key performance indicators for each model. It can be seen that the EfficientNet transfer learning model and the main model have very comparable performance - with the EfficientNet model and main model achieving an accuracy




Fig. 4. Confusion Matrix of main model and transfer learning model

of 85.6% and 85.4% respectively - a difference of only 0.2%. Furthermore the EfficientNet and main model achieved F1 scores of 0.862 and 0.868 respectively, a difference of only 0.006.

Where these models differ is their precision and recall. The EfficientNet model achieved a precision of 87.5% (4.2% higher than the main models 83.3%.) The main model achieved a recall of 90.6% (5.7% higher than the EfficientNet models 84.9%.) This means that the EfficientNet model is better at minimising false positive AI detection, whereas the main model is better at detecting AI cases without missing them (i.e. reducing false negatives).

This is evident in fig. 4, where the EfficientNet model has 32 false posiives compared to the main models 48. Similarly, the main model has only 25 false negatives, compared to the EfficientNet models 40.

AWS Sagemaker Information and Discussion

This project was developed using AWS Sagemaker. Figure 5, which overall was less than \$15. This was achieved due to careful management of resources with a known budget of \$50 leading to selection of smaller instance types (using general purpose ml.c5.2xlarge instances for training and inference rather than larger, more expensive GPU instances). This of course, came at the cost of longer training times, but was deemed acceptable given the budget constraints, and the fact that the workflow would be very easily extensible to a more expensive configuration in a real world deployment scenario.

mention the drawbacks of the deployment here, though have the code chunk (with the creation of the endpoint and the prediction extraction, (and in the appendix have the code snippets for the table and the confusion matrix))

Final Model Discussion and Conclusion

This report proposed two models to detect if an image is generated using ai - one model developed through hyperparameter tuning, and the other through transfer learning from EfficientNet. The results indicate that both models perform comparably, with slight differences in precision and recall.

The original research objective was to make a tool to distinguish between AI-generated and real images, ultimately as a means to identify the spread of misinformation and other harms in online discourse. Were this model to be deployed as a 'first step' with the intention of more thorough investigation, it would

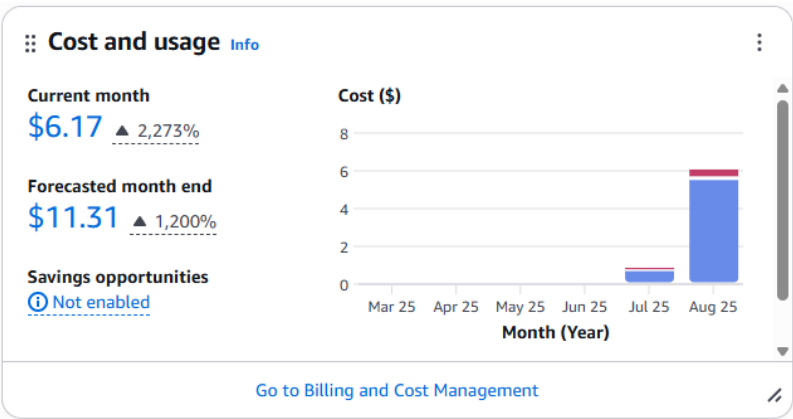


Fig. 5. Project Cost summary

be more beneficial to be relaxed on false positives, and minimise false negatives, given the false positives would be vindicated with further analysis or research.

It is for this reason that, assuming this approach, the main model would be the preferred model, as it has a higher recall (90.6% vs 84.9%) and therefore is slightly better at detecting AI-generated images without missing them. talk about the paper at the start and that we didnt achieve research level performance, nor the experience to use the dataset they made

mention the limitation of fixed resolutions/aspect ratio of 1:1

Overall it can be concluded that the models developed in this investigation are effective at distinguishing AI-generated images from other images, successfully training and deploying 2 models to make this prediction.

References

- [1] Z. Lu, D. Huang, L. Bai, J. Qu, C. Wu, X. Liu, and W. Ouyang, “Seeing is not always believing: Benchmarking human and model perception of ai-generated images,” 2023. [Online]. Available: <https://arxiv.org/abs/2304.13023>
- [2] C. Scott *et al.* (2024) Exploring the ethics of artificial intelligence in art. Accessed: 22 July 2025. [Online]. Available: <https://www.artshub.com.au/news/opinions-analysis/exploring-the-ethics-of-artificial-intelligence-in-art-2694121/>
- [3] J. J. Bird and A. Lotfi, “Cifake: Image classification and explainable identification of ai-generated synthetic images,” 2023. [Online]. Available: <https://arxiv.org/abs/2303.14126>
- [4] Hugging Face Datasets, “competitions/aionot,” <https://huggingface.co/datasets/competitions/aionot>, 2023, accessed: 22 July 2025.
- [5] M. A. K. Raiaan, S. Sakib, N. M. Fahad, A. A. Mamun, M. A. Rahman, S. Shatabda, and M. S. H. Mukta, “A systematic review of hyperparameter optimization techniques in convolutional neural networks,” *Decision Analytics Journal*, vol. 11, p. 100470, 2024. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S2772662224000742>
- [6] K. Team. (2024) Keras applications. Accessed: 22 July 2025. [Online]. Available: <https://keras.io/api/applications/>
- [7] M. Tan and Q. V. Le, “Efficientnetv2: Smaller models and faster training,” in *International Conference on Machine Learning*, 2021, iCML 2021. [Online]. Available: <https://arxiv.org/abs/2104.00298>

Appendix

Listing 2. Data Wrangling Script

```

1
2 with open("token.txt", "r") as file:
3     hugging_face_token = file.read().strip() # must have credentials to access the dataset
4
5 raw_dataset = load_dataset('competitions/aiornot')
6 raw_dataset = raw_dataset['train'] # remove the unused 'test' set with no labels
7 dataset_length = len(raw_dataset)
8
9 # first holdout 500 images for final model evaluation
10 holdout_test = raw_dataset.select(range(500))
11 dataset = raw_dataset.select(range(500, len(raw_dataset)))
12
13 # split dataset into:
14 # - small_train (10% of original)
15 # - train (70% of original)
16 # - test (20% of original)
17
18 # first split: small_train (10%) and remainder (90%)
19 split_1 = dataset.train_test_split(train_size=0.1, seed=RANDOM_SEED)
20 small_train = split_1["train"]
21 remainder = split_1["test"]
22
23 # second split: train (70%) and test (20%)
24 split_2 = remainder.train_test_split(train_size=0.7 / 0.9, seed=RANDOM_SEED)
25 train = split_2["train"]
26 test = split_2["test"]
27
28 is_any_data_unused = not(holdout_test.num_rows + small_train.num_rows + train.num_rows + test.num_rows == dataset.num_rows)
29 assert not is_any_data_unused, "Some data is unused in the splits."
30
31 def to_numpy(dataset):
32     images = np.stack([np.asarray(img) for img in dataset["image"]])
33     labels = np.array(dataset["label"])
34     return images, labels
35
36 x_small, y_small = to_numpy(small_train)
37 x_train, y_train = to_numpy(train)
38 x_test, y_test = to_numpy(test)
39 x_holdout, y_holdout = to_numpy(holdout_test)
40
41 sampler = RandomUnderSampler(random_state=RANDOM_SEED)
42
43 def resample_images_and_labels(images, labels, sampler):
44     flat_images = images.reshape((images.shape[0], -1))
45     resampled_flat, resampled_labels = sampler.fit_resample(flat_images, labels)
46     resampled_images = resampled_flat.reshape((-1,) + images.shape[1:])
47     return resampled_images, resampled_labels
48
49 x_small_resampled, y_small_resampled = resample_images_and_labels(x_small, y_small, sampler)
50 x_train_resampled, y_train_resampled = resample_images_and_labels(x_train, y_train, sampler)
51 x_test_resampled, y_test_resampled = resample_images_and_labels(x_test, y_test, sampler)
52
53 # save to npz file and upload to s3 bucket

```